

Generic Machine Learning Inference on Heterogeneous Treatment Effects Using the Package GenericML

Max Welz^{1,2} Andreas Alfons¹ Mert Demirer³
Victor Chernozhukov³

¹Erasmus School of Economics, Erasmus University Rotterdam

²Erasmus University Medical Center (Erasmus MC)

³Massachusetts Institute of Technology

useR!, June 21, 2022



Motivation

Recent literature in causal inference is focused on heterogeneous treatment effects

- Often based on Machine Learning (ML) techniques
- Goal: Consistent estimation and uniformly valid inference on conditional average treatment effect (CATE)

→ Difficult w/o strong assumptions, especially in high dimensions!

→ Generic Machine Learning Inference (Generic ML; Chernozhukov, Demirer, Duflo, and Fernández-Val, 2020) remedies this in randomized experiments

Motivation

Recent literature in causal inference is focused on heterogeneous treatment effects

- Often based on Machine Learning (ML) techniques
- Goal: Consistent estimation and uniformly valid inference on conditional average treatment effect (CATE)

→ Difficult w/o strong assumptions, especially in high dimensions!

→ Generic Machine Learning Inference (Generic ML; Chernozhukov, Demirer, Duflo, and Fernández-Val, 2020) remedies this in randomized experiments

Motivation

Recent literature in causal inference is focused on heterogeneous treatment effects

- Often based on Machine Learning (ML) techniques
- Goal: Consistent estimation and uniformly valid inference on conditional average treatment effect (CATE)

→ Difficult w/o strong assumptions, especially in high dimensions!

→ **Generic Machine Learning Inference** (Generic ML; Chernozhukov, Demirer, Duflo, and Fernández-Val, 2020) remedies this in randomized experiments

Setup

Let

- Y be the outcome
- Z be a possibly high-dimensional vector of covariates
- D be a binary treatment assignment variable

→ Observe $(Y_i, Z_i, D_i)_{i=1}^N$ as i.i.d. copies of (Y, Z, D)

→ Assume unconfoundedness and random treatment assignment

A Very General Model

We consider the very general model

$$Y = b_0(Z) + Ds_0(Z) + U, \quad E[U \mid Z, D] = 0,$$

where

$$b_0(Z) = E[Y \mid D = 0, Z]$$

is the baseline conditional average (BCA), and

$$s_0(Z) = E[Y \mid D = 1, Z] - E[Y \mid D = 0, Z]$$

is the conditional average treatment effect (CATE)

A Very General Model

We consider the very general model

$$Y = b_0(Z) + Ds_0(Z) + U, \quad E[U \mid Z, D] = 0,$$

where

$$b_0(Z) = E[Y \mid D = 0, Z]$$

is the **baseline conditional average (BCA)**, and

$$s_0(Z) = E[Y \mid D = 1, Z] - E[Y \mid D = 0, Z]$$

is the conditional average treatment effect (CATE)

A Very General Model

We consider the very general model

$$Y = b_0(Z) + D s_0(Z) + U, \quad E[U \mid Z, D] = 0,$$

where

$$b_0(Z) = E[Y \mid D = 0, Z]$$

is the baseline conditional average (BCA), and

$$s_0(Z) = E[Y \mid D = 1, Z] - E[Y \mid D = 0, Z]$$

is the conditional average treatment effect (CATE)

Focus of Generic ML

Generic ML focuses on estimation and inference on

key features of $s_0(Z)$ rather than $s_0(Z)$ itself

→ No need for consistent estimation of $s_0(Z)$ or $b_0(Z)$!

Focus of Generic ML

Generic ML focuses on estimation and inference on

key features of $s_0(Z)$ rather than $s_0(Z)$ itself

→ No need for consistent estimation of $s_0(Z)$ or $b_0(Z)$!

Generic ML

- ➊ Randomly partition the data in two disjoint sets A and M
- ➋ On set A , use some machine learner to obtain estimates $B(Z)$ and $S(Z)$ of $b_0(Z)$ and $s_0(Z)$, respectively
- ➌ On set M , calculate the key features of $s_0(Z)$

Two sources of uncertainty:

- Estimation uncertainty (conditional on set A) from Step 2
- Splitting uncertainty from the sample splitting in Step 1

→ Address by repeating Steps 1–3 many times

A stylized, handwritten-style logo for Erasmus, featuring a large, flowing 'E' followed by the word 'Erasmus' in a cursive script.

Generic ML

- ➊ Randomly partition the data in two disjoint sets A and M
- ➋ On set A , use some machine learner to obtain estimates $B(Z)$ and $S(Z)$ of $b_0(Z)$ and $s_0(Z)$, respectively
- ➌ On set M , calculate the key features of $s_0(Z)$

Two sources of uncertainty:

- Estimation uncertainty (conditional on set A) from Step 2
- Splitting uncertainty from the sample splitting in Step 1

→ Address by repeating Steps 1–3 many times

A stylized, handwritten-style logo for Erasmus, featuring a large, flowing 'E' followed by the word 'Erasmus' in a cursive script.

Generic ML

- ➊ Randomly partition the data in two disjoint sets A and M
- ➋ On set A , use some machine learner to obtain estimates $B(Z)$ and $S(Z)$ of $b_0(Z)$ and $s_0(Z)$, respectively
- ➌ On set M , calculate the key features of $s_0(Z)$

Two sources of uncertainty:

- Estimation uncertainty (conditional on set A) from Step 2
- Splitting uncertainty from the sample splitting in Step 1

→ Address by repeating Steps 1–3 many times

A stylized, handwritten-style logo of the word "Erasmus" in a dark blue or black color.

Inference

Variational Estimation and Inference (VEIN):

- Fix significance level $\alpha \in (0, 0.5)$
 - Calculate the key features across S splits of the data
 - Take medians across the S splits of each key feature parameter
- Inference on each key feature parameter with size control of level 2α
- Can be repeated for many machine learners (report the “best” one)

Software Implementation

Package **GenericML** (Welz, Alfons, Demirer, and Chernozhukov, 2022)

- CRAN: <https://cran.r-project.org/package=GenericML>
- GitHub: <https://github.com/mwelz/GenericML>

→ Flexible, user-friendly, fast, object-oriented

→ Based on mlr3 ecosystem of Lang et al. (2019)

Empirical Example: Setup

We revisit Crépon et al.'s (2015) study on the effects of microcredits¹

- Sample: 162 villages in rural Morocco, divided into 81 similar pairs
- Randomly select one village in each pair and make microcredits available for the residents
- Measure if total borrowing changes

¹We thank Esther Duflo for making the data available to us

Empirical Example: Data

Household-level data on $N = 5,513$ households

- **Dependent variable Y :** total volume of borrowing
- **Treatment indicator D :** 1 if household can access microcredits
- **Covariates Z :** 97 variables (after encoding), among which
 - `head_age_b1` is age of household's head
- **Grouping variables:**
 - `demi_paire` is a factor of village membership
 - `vil_pair` is a factor of village pair membership

Empirical Example: Baseline Results

Crépon et al. (2015) find that microcredit availability has...

- low take-up (17% in treatment group)
- significant effect on total borrowing: ATE of MAD² 1,206 ($p < 0.01$)

→ Use GenericML to investigate heterogeneity in this effect!

²MAD = Moroccan Dirham

Empirical Example: Specification of Learners

- Specify a suite of learners with `mlr3` syntax
- Here: random forest, elastic net, support vector machine, gradient boosting

```
R> # devtools::install_github("mwelz/GenericML")
R> # version 0.2.3, not yet on CRAN!
R> library("GenericML")
R>
R> # load data, available in GitHub repo mwelz/GenericML
R> load("slides/data/morocco_preprocessed.Rdata")
R>
R> # specify learners
R> learners <-
+   c("random_forest",
+     "mlr3::lrn('cv_glmnet', s = 'lambda.min', alpha = 0.5)",
+     "mlr3::lrn('svm')",
+     "mlr3::lrn('xgboost')")
```

Empirical Example: Specification of Learners

- Specify a suite of learners with `mlr3` syntax
- Here: random forest, elastic net, support vector machine, gradient boosting

```
R> # devtools::install_github("mwelz/GenericML")
R> # version 0.2.3, not yet on CRAN!
R> library("GenericML")
R>
R> # load data, available in GitHub repo mwelz/GenericML
R> load("slides/data/morocco_preprocessed.Rdata")
R>
R> # specify learners
R> learners <-
+   c("random_forest",
+     "mlr3::lrn('cv_glmnet', s = 'lambda.min', alpha = 0.5)",
+     "mlr3::lrn('svm')",
+     "mlr3::lrn('xgboost')")
```

Empirical Example: Customization

Spatial data of 81 village pairs

- Include fixed effects for each pair
- Cluster standard errors on the village level
- `GenericML` allows this through setup functions
- Support for sandwich covariance estimators (Zeileis, 2004)

Empirical Example: Customization

`setup_X1()` customizes inclusion of controls and fixed effects

```
R> # include BCA and CATE controls
R> # add fixed effects along variable "vil_pair"
R> X1 <- setup_X1(funs_Z = c("B", "S"),
+               fixed_effects = vil_pair)
```

`setup_vcov()` customizes covariance estimation

```
R> # calls functions from the "sandwich" package
R> # cluster standard errors along "demi_paire"
R> vcov <- setup_vcov(estimator = "vcovCL",
+                   arguments = list(cluster = demi_paire))
```

Empirical Example: Customization

`setup_X1()` customizes inclusion of controls and fixed effects

```
R> # include BCA and CATE controls
R> # add fixed effects along variable "vil_pair"
R> X1 <- setup_X1(funs_Z = c("B", "S"),
+               fixed_effects = vil_pair)
```

`setup_vcov()` customizes covariance estimation

```
R> # calls functions from the "sandwich" package
R> # cluster standard errors along "demi_paire"
R> vcov <- setup_vcov(estimator = "vcovCL",
+                   arguments = list(cluster = demi_paire))
```

Empirical Example: Customization

`setup_X1()` customizes inclusion of controls and fixed effects

```
R> # include BCA and CATE controls
R> # add fixed effects along variable "vil_pair"
R> X1 <- setup_X1(funs_Z = c("B", "S"),
+               fixed_effects = vil_pair)
```

`setup_vcov()` customizes covariance estimation

```
R> # calls functions from the "sandwich" package
R> # cluster standard errors along "demi_paire"
R> vcov <- setup_vcov(estimator = "vcovCL",
+                   arguments = list(cluster = demi_paire))
```


Empirical Example: Customization

`setup_X1()` customizes inclusion of controls and fixed effects

```
R> # include BCA and CATE controls
R> # add fixed effects along variable "vil_pair"
R> X1 <- setup_X1(funs_Z = c("B", "S"),
+               fixed_effects = vil_pair)
```

`setup_vcov()` customizes covariance estimation

```
R> # calls functions from the "sandwich" package
R> # cluster standard errors along "demi_paire"
R> vcov <- setup_vcov(estimator = "vcovCL",
+                   arguments = list(cluster = demi_paire))
```

A stylized, handwritten-style logo of the word "Erasmus" in a dark grey or black color.

GenericML Interface

```
R> x <- GenericML(  
+   Z = Z, D = D, Y = Y,           # observed data  
+   learners_GenericML = learners, # learners  
+   learner_propensity_score = "constant", # = 0.5 (RCT)  
+   num_splits = 100L,             # number splits  
+   quantile_cutoffs = c(0.2, 0.4, 0.6, 0.8), # grouping  
+   significance_level = 0.05,      # significance level  
+   X1_BLP = X1, X1_GATES = X1,    # regression setup  
+   vcov_BLP = vcov, vcov_GATES = vcov, # covariance setup  
+   parallel = TRUE, num_cores = 6L, # parallelization  
+   seed = 20220621)               # RNG seed
```

...and many more arguments for fine-tuning!

→ stratified sampling, Horvitz-Thompson transformation...



GenericML Interface

```
R> x <- GenericML(  
+   Z = Z, D = D, Y = Y,           # observed data  
+   learners_GenericML = learners, # learners  
+   learner_propensity_score = "constant", # = 0.5 (RCT)  
+   num_splits = 100L,             # number splits  
+   quantile_cutoffs = c(0.2, 0.4, 0.6, 0.8), # grouping  
+   significance_level = 0.05,      # significance level  
+   X1_BLP = X1, X1_GATES = X1,     # regression setup  
+   vcov_BLP = vcov, vcov_GATES = vcov, # covariance setup  
+   parallel = TRUE, num_cores = 6L, # parallelization  
+   seed = 20220621)               # RNG seed
```

...and many more arguments for fine-tuning!

→ stratified sampling, Horvitz-Thompson transformation...



Analysis of GenericML Objects

Methods for the analysis of the *key features* of CATE

- `get_BLP()`
- `get_GATES()`
- `get_CLAN()`

→ linked to `rich plot()` and `print()` methods

Empirical Example: get_BLP()

Best Linear Predictor (BLP): Estimates some (β_1, β_2) via OLS:

- $\beta_1 = E s_0(Z)$ is the ATE
- $\beta_2 \neq 0$ if there is heterogeneity in $s_0(Z)$ and $S(Z)$ predicts it well

```
R> get_BLP <- get_BLP(x, plot = TRUE)
R> get_BLP # print method
BLP generic targets
---
      Estimate    CI lower CI upper p value
beta.1 1113.50155  273.02645 1935.274 0.00945 **
beta.2   0.35315  -0.04384   0.698 0.08613 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
---
Confidence level of confidence interval [CI lower, CI upper]: 90 %
```

Empirical Example: get_BLP()

Best Linear Predictor (BLP): Estimates some (β_1, β_2) via OLS:

- $\beta_1 = E s_0(Z)$ is the ATE
- $\beta_2 \neq 0$ if there is heterogeneity in $s_0(Z)$ and $S(Z)$ predicts it well

```
R> get_BLP <- get_BLP(x, plot = TRUE)
```

```
R> get_BLP # print method
```

```
BLP generic targets
```

```
----
```

	Estimate	CI lower	CI upper	p value	
beta.1	1113.50155	273.02645	1935.274	0.00945	**
beta.2	0.35315	-0.04384	0.698	0.08613	.

```
----
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
----
```

```
Confidence level of confidence interval [CI lower, CI upper]: 90 %
```

Empirical Example: get_BLP()

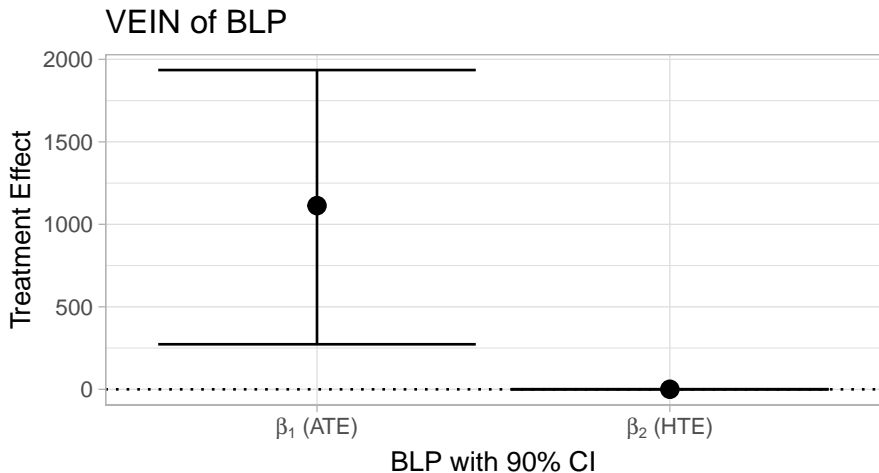
Best Linear Predictor (BLP): Estimates some (β_1, β_2) via OLS:

- $\beta_1 = E s_0(Z)$ is the ATE
- $\beta_2 \neq 0$ if there is heterogeneity in $s_0(Z)$ and $S(Z)$ predicts it well

```
R> get_BLP <- get_BLP(x, plot = TRUE)
R> get_BLP # print method
BLP generic targets
---
      Estimate    CI lower CI upper p value
beta.1 1113.50155  273.02645 1935.274 0.00945 **
beta.2   0.35315  -0.04384   0.698 0.08613 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
---
Confidence level of confidence interval [CI lower, CI upper]: 90 %
```

Empirical Example: get_BLP()

```
R> plot(get_BLP) # plot method
```



Empirical Example: `get_GATES()`

Sorted Group Average Treatment Effects (GATES): Build groups

$$G_k := \{S(Z) \in I_k\}, \quad k = 1, \dots, K,$$

where $I_k = [\ell_{k-1}, \ell_k)$ divide the support of $S(Z)$ into regions

→ Estimate group-ATE $\gamma_k := E[s_0(Z) \mid G_k]$ via OLS

Empirical Example: get_GATES()

```
R> get_GATES <- get_GATES(x, plot = TRUE)
R> get_GATES
GATES generic targets
---
```

	Estimate	CI lower	CI upper	p value
gamma.1	-80.44	-2517.30	2097	0.93525
gamma.2	305.50	-674.10	1336	0.49251
gamma.3	725.63	-505.53	1932	0.19349
gamma.4	1744.51	395.93	3097	0.01225 *
gamma.5	2743.76	759.85	4940	0.00911 **
gamma.5-gamma.1	2922.13	-89.43	6087	0.05536 .

```
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

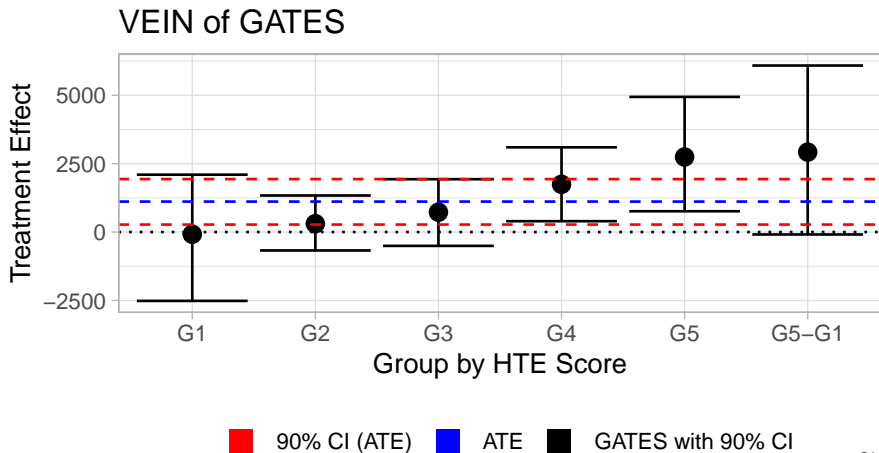
```
---
```

Confidence level of confidence interval [CI lower, CI upper]: 90 %



Empirical Example: get_GATES()

```
R> plot(get_GATES)
```



Empirical Example: get_CLAN()

Classification Analysis (CLAN): Observed within-group averages, δ_k , of a variable for groups G_k

```
R> get_CLAN <- get_CLAN(x, variable = "head_age_bl", plot = TRUE)
R> get_CLAN
```

CLAN generic targets for variable 'head_age_bl'

```
---
              Estimate CI lower CI upper  p value
delta.1          36.49   34.46   38.554 < 2e-16 ***
delta.2          43.66   42.12   45.210 < 2e-16 ***
delta.3          41.40   39.50   43.258 < 2e-16 ***
delta.4          34.75   32.55   36.853 < 2e-16 ***
delta.5          23.85   21.53   26.151 < 2e-16 ***
delta.5-delta.1  -12.52  -15.61   -9.514 4.44e-16 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
---
```

```
Confidence level of confidence interval [CI lower, CI upper]: 90 %
```

Empirical Example: get_CLAN()

Classification Analysis (CLAN): Observed within-group averages, δ_k , of a variable for groups G_k

```
R> get_CLAN <- get_CLAN(x, variable = "head_age_bl", plot = TRUE)
R> get_CLAN
```

CLAN generic targets for variable 'head_age_bl'

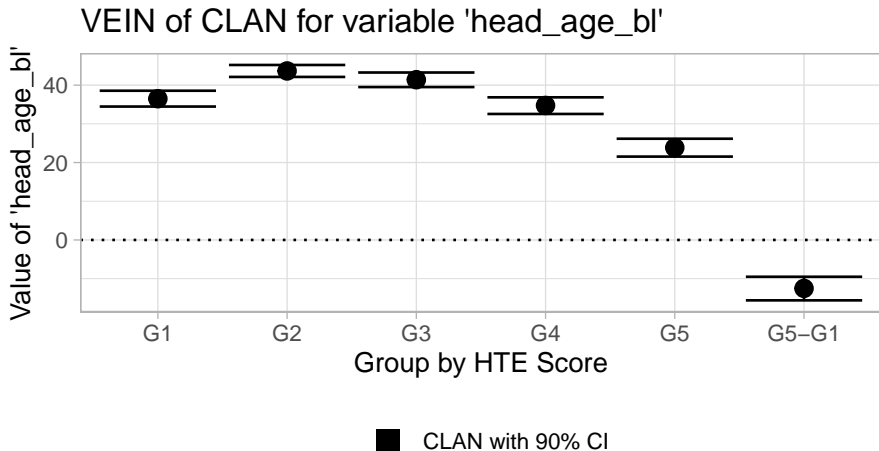
```
---
              Estimate CI lower CI upper  p value
delta.1          36.49   34.46   38.554 < 2e-16 ***
delta.2          43.66   42.12   45.210 < 2e-16 ***
delta.3          41.40   39.50   43.258 < 2e-16 ***
delta.4          34.75   32.55   36.853 < 2e-16 ***
delta.5          23.85   21.53   26.151 < 2e-16 ***
delta.5-delta.1  -12.52  -15.61   -9.514 4.44e-16 ***
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
---
Confidence level of confidence interval [CI lower, CI upper]: 90 %
```

Empirical Example: get_CLAN()

```
R> plot(get_CLAN)
```



Conclusions and Discussion

Conclusions

- High-dimensional uniformly valid inference on CATE is hard
- Generic ML can do so under minimal assumptions by focusing on key features of CATE instead of CATE itself
- R package `GenericML` available on CRAN

Future work

- Implement monotonization of confidence bounds
- Enable support for deep learning, perhaps via `mlr3keras`

Conclusions and Discussion

Conclusions

- High-dimensional uniformly valid inference on CATE is hard
- Generic ML can do so under minimal assumptions by focusing on key features of CATE instead of CATE itself
- R package `GenericML` available on CRAN

Future work

- Implement monotonization of confidence bounds
- Enable support for deep learning, perhaps via `mlr3keras`

References

- Victor Chernozhukov, Mert Demirer, Esther Duflo, and Iván Fernández-Val. Generic Machine Learning Inference on Heterogenous Treatment Effects in Randomized Experiments. **arXiv preprint: arXiv:1712.04802**, 2020.
- Bruno Crépon, Florencia Devoto, Esther Duflo, and William Parienté. Estimating the Impact of Microcredit on Those Who Take It Up: Evidence from a Randomized Experiment in Morocco. **American Economic Journal: Applied Economics**, 7(1):123–150, 2015.
- Michel Lang, Martin Binder, Jakob Richter, Patrick Schratz, Florian Pfisterer, Stefan Coors, Quay Au, Giuseppe Casalicchio, Lars Kotthoff, and Bernd Bischl. `mlr3`: A Modern Object-Oriented Machine Learning Framework in R. **Journal of Open Source Software**, 4(44):1903, 2019.
- Max Welz, Andreas Alfons, Mert Demirer, and Victor Chernozhukov. **GenericML: Generic Machine Learning Inference**, 2022. URL <https://CRAN.R-project.org/package=GenericML>. R package version 0.2.2.
- Achim Zeileis. Econometric Computing with HC and HAC Covariance Matrix Estimators. **Journal of Statistical Software**, 11(10):1–17, 2004.



Algorithm 1 in Chernozhukov et al. (2020)

IN: Data = $(Y_i, Z_i, D_i)_{i=1}^N$, significance level α , a suite of ML methods, number of splits S

OUT: p -values and $(1 - 2\alpha)$ confidence intervals of point estimates of each target parameter in GATES, BLP, and CLAN

- ① Compute propensity scores $p(Z_i), i = 1, \dots, N$
- ② Do S splits of $\{1, \dots, N\}$ into disjoint sets A and M of same size
- ③ **for** each ML method and each split $s = 1, \dots, S$, **do**
 - Ⓐ Tune and train each ML method to learn $B(\cdot)$ and $S(\cdot)$ on A
 - Ⓑ On M , use $B(\cdot)$ and $S(\cdot)$ to estimate the BLP, GATES, CLAN target parameters
 - Ⓒ Compute some performance measures for the ML methods
- ④ Choose the best ML method based on the medians of the performance measures
- ⑤ Calculate the medians of the confidence bounds, p -values, and point estimates of each target parameter
- ⑥ Adjust the confidence bounds and p -values



Best Learner

Compute two performance measures for each learner

$$\hat{\Lambda} = |\hat{\beta}_2|^2 \widehat{\text{Var}}(S(Z)), \quad \hat{\bar{\Lambda}} = \frac{1}{K} \sum_{k=1}^K \hat{\gamma}_k^2$$

→ Best learner maximizes their median across S splits

→ In the empirical example, that's random forest (get via `get_best()`)